



AgentForge

A mobile-first agentic AI platform with 247 specialist agents, multi-agent pipelines, and live AI debate — running 100% in the browser. No server. No build step. No cost.

247

AI SPECIALISTS

16

DIVISIONS

6

INTERACTION MODES

5

AI PROVIDERS




mmuzammul.github.io/AgForge

Project Overview & Capability Brief

EXECUTIVE SUMMARY

One app. 247 AI experts. Zero installs.

AgentForge turns a single static web page into a full agentic AI workforce. Instead of one generic chatbot, every conversation is led by a purpose-built specialist — a Backend Architect, a Growth Hacker, a Security Auditor — each with a deep, hand-written system prompt. Agents can work alone, in a chained pipeline, as an auto-assembled team, or even argue with each other to pressure-test an answer. Everything runs client-side: no backend, no database, no per-seat licensing, and no vendor lock-in.



WHY IT MATTERS

- ▶ **Speed to value** — open the link, pick a specialist, get expert-level output in seconds. No onboarding, no account, no setup.
- ▶ **Built for real work** — Pipeline and Auto-Build Team modes chain multiple specialists so a single goal (e.g. "launch a SaaS landing page") produces a coordinated, multi-discipline output.
- ▶ **Free-tier friendly** — a 5-provider failover chain (Groq → Cerebras → Gemini → OpenRouter → Ollama) keeps long pipelines running even on free API quotas.
- ▶ **Honest engineering** — no frameworks, no bundler, no hidden backend. The entire product is auditable in a handful of plain JavaScript files.

Why AgentForge, not just “another chatbot”

Most AI chat tools give you one general-purpose model and a blank text box. AgentForge gives you a curated bench of 247 experts and the orchestration tooling to put them to work together.

CAPABILITY	AGENTFORGE	GENERIC AI CHAT
247 domain specialists, each a deep expert system prompt	✓	✗
Auto-builds the optimal team for your stated goal	✓	✗
Multi-agent pipelines with chained reasoning context	✓	✗
Agent-vs-agent debate to stress-test an answer	✓	✗
LLM-judged benchmarker across multiple agents	✓	✗
Runs the code it writes, 70+ languages, one tap	✓	✗
Voice input · file attachment · memory · export	✓	sometimes
Zero install, zero dependencies, free forever	✓	✗
Installable as a native-feeling app (PWA, offline shell)	✓	✗
Your data never leaves your browser	✓	✗

THE ROSTER

247 agents across 16 professional divisions

Every agent is a hand-written Markdown system prompt — version-controlled, auditable, and fetched on demand. The roster spans full software teams, growth & marketing, finance, security, GIS, game development, and more.



CAPABILITIES

Six ways to put the roster to work

One specialist isn't always enough. AgentForge offers six distinct interaction modes so the right tool fits the job — from a quick question to a fully orchestrated, multi-agent deliverable.



01

Chat

Talk to any of the 247 specialists one-on-one, with streaming responses, file attachments, voice input, and full conversation history saved per agent.



02

Agent Pipeline

Chain 2–6 specialists so each one builds on the previous agent's output. Run a full product launch — Product Manager → Frontend Developer → Brand Guardian → Growth Hacker — in a single tap.



03

Auto-Build Team

Describe a goal in plain English; an orchestrator reads the entire 247-agent roster and assembles the right specialists, in the right order, automatically.



04

Agent Debate

Two specialists argue in configurable rounds — a Proposer builds the case, a Critic challenges it — until a battle-tested, production-ready answer emerges.



05

Benchmarker

Run the same prompt across multiple agents at once; an impartial judge model scores every response on relevance, depth, and correctness.



06

Code Runner

Every code block produced by any agent gets a Run button — executed live via the free Piston API, covering 70+ programming languages.

LIVE APPLICATION

See it in action

Real screenshots captured directly from the running application — not mockups.

AgentForge

Search 247 specialists...

Allacademic 13design 16engineering 43finance 10

POPULAR

**Backend Architect**
engineering
Senior backend architect specializing in scalable system design, database architecture, API developm

**Frontend Developer**
engineering
Expert frontend developer specializing in modern web technologies, React/Vue/Angular frameworks, UI

**AI Engineer**
engineering
Expert AI/ML engineer specializing in machine learning model development, deployment, and...

**Growth Hacker**
marketing
Expert growth strategist specializing in rapid user acquisition through data-driven experimentation.

**Product Manager**
product
Holistic product leader who owns the full product lifecycle — from discovery and strategy through ro

**UI Designer**
design
Expert UI designer specializing in visual design systems, component libraries, and pixel-perfect int

**SEO Specialist**
marketing
Expert search engine optimization strategist specializing in technical SEO, content optimization, li

**Penetration Tester**
security
Offensive security specialist conducting authorized penetration tests, red team operations, and vuln

**Business Strategist**
specialized

Agent roster & search

AgentForge

Backend Architect
engineering

How would you design a backend for a real-time chat app with 1M users?

CopyQuote

**Hi! I'm Backend Architect**
Designs the systems that hold everything up — databases, APIs, cloud, scale.
Tap  → connect free AI for real responses.

2.4s
CopyQuoteGist

 Message...

1:1 chat with a specialist

AgentForge

PROJECT GOAL

Describe what you want to build or solve...
e.g. Build a SaaS landing page for a project management tool

Auto-Build Team

AGENT PIPELINE (IN ORDER)

1**Product Manager**
product

2**Frontend Developer**
engineering

3**Brand Guardian**
design

4**Growth Hacker**
marketing

Run Pipeline

Share Pipeline

EXAMPLE PIPELINES

**SaaS Product Launch**
Product Manager → Frontend Dev → Brand Guardian → Marketing

**Security Audit**
Security Architect → AppSec Engineer → QA Tester

**Content Campaign**
Brand Guardian → Content Strategist → SEO Specialist

Pipeline queue (4 agents)

AgentForge

Demo

✕ AGENT VS AGENT DEBATE

Two specialists argue and refine until a production-ready answer emerges.

AGENT A — PROPOSER

Backend Architect · engineering

AGENT B — CRITIC

Security Architect · security

TOPIC / QUESTION

Should we use microservices or a monolith for a 10-person startup?

ROUNDS

234

Start Debate

Agent vs. Agent debate setup

AgentForge

Demo

AGENT BENCHMARKER

Run the same prompt across agents and rank by quality, depth & speed.

Design a REST API for user authentication with JWT

AGENTS TO COMPARE (SELECT 2+)

☒ Frontend Developer
engineering

☒ Backend Architect
engineering

☒ Software Architect
engineering

☒ UI Designer
design

Run Benchmark

Benchmarker, 4 agents selected

AgentForge

Demo

Saved Memory

Clear all

Backend Architect

Designed a horizontally-scalable backend using event-driven microservices, Redis pub/sub, and...

Today 10:42 AM

chat

ViewCopyDelete

Pipeline: SaaS landing page launch

Full launch plan covering product positioning, frontend build, brand guidelines, and growth loops...

Today 9:15 AM

pipeline

ViewCopyDelete

Debate: Microservices vs Monolith

After 3 rounds, both agents converged on a modular monolith with clear service boundaries...

Yesterday

chat

ViewCopyDelete

Saved memory / outputs

ENGINEERING

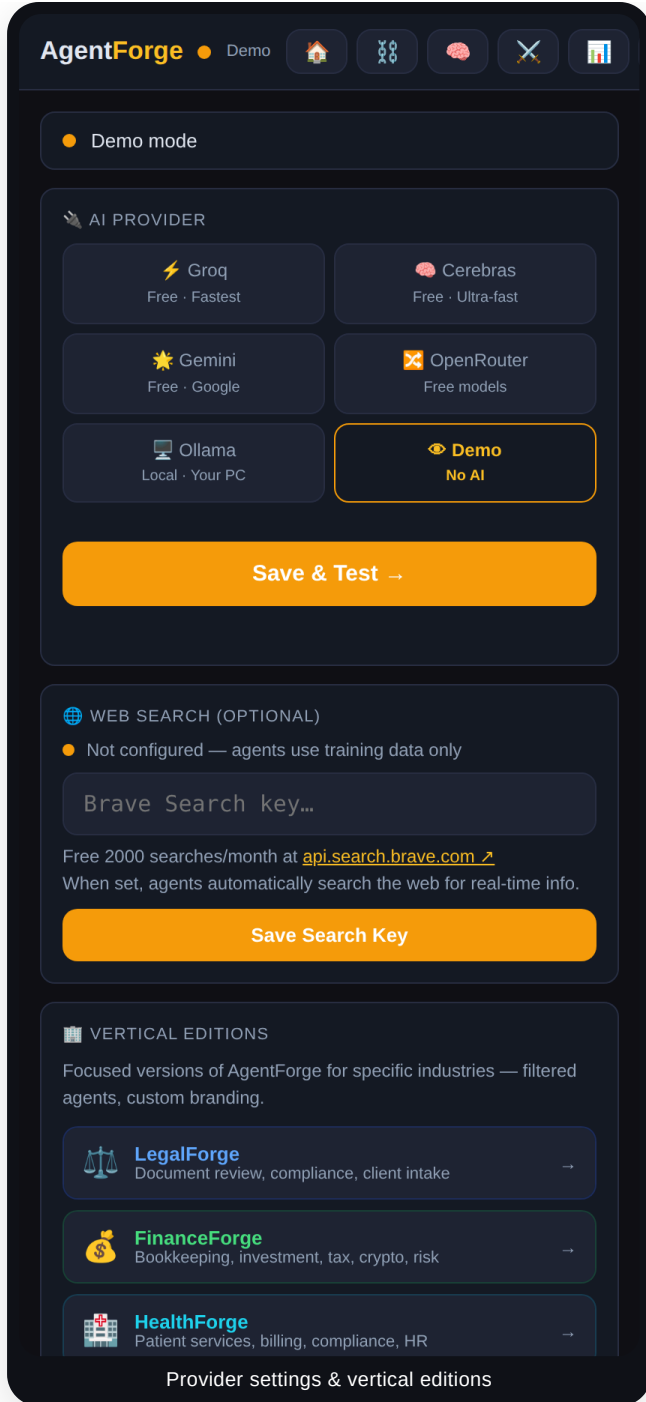
Built simple, on purpose

AgentForge is deliberately dependency-free: static files, no build step, no bundler. The entire product is a handful of plain JavaScript files sharing one global scope — fully auditable, fast to load on mobile, and trivial to self-host.

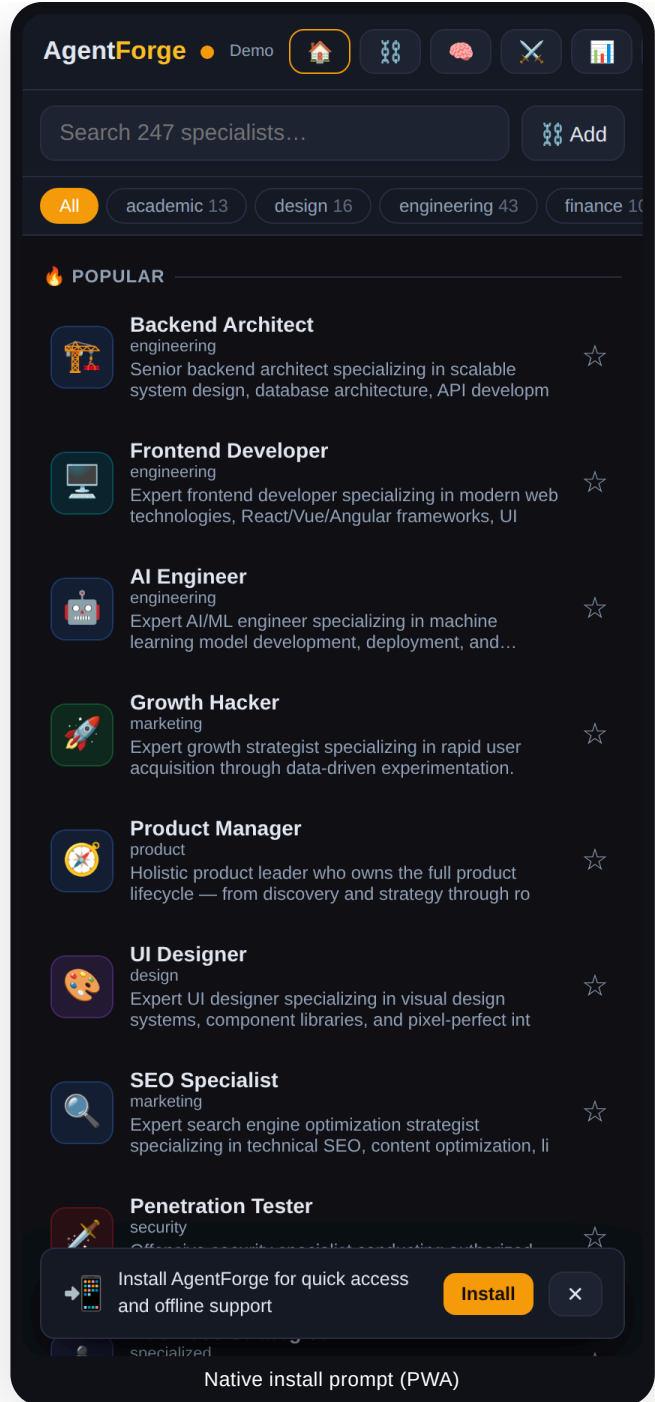
```
index.html  (markup only)
assets/css/styles.css  (all styling, CSS variables)
assets/js/ data.js · ui.js · chat.js · pipeline.js · providers.js · memory.js ·
settings.js · modes.js · main.js
  → fetches agents/<division>/<id>.md on demand from this repo
  → streams responses from Groq → Cerebras → Gemini → OpenRouter → Ollama
```

SMART FAILOVER, BY DESIGN

- ▶ A 5-provider cascade keeps multi-agent pipelines running unattended, even on free-tier rate limits.
- ▶ On a 429, the engine parses the suggested wait time, rotates models, and retries automatically.
- ▶ Agent system prompts are fetched lazily and cached, keeping the initial page load fast on mobile data.
- ▶ Installable as a Progressive Web App with offline shell support, on both Android and iOS.
- ▶ All state — settings, chats, pipelines, saved memory — lives in the browser's `localStorage`. Nothing is sent to a backend this project owns, because there is no backend.



AgentForge — Project Overview



Native install prompt (PWA)

THE FULL ROSTER

All 16 divisions, at a glance

Engineering	43	Specialized	45
Marketing	32	Design	16
Academic	13	Security	12
Sales	12	Finance	10
GIS & Spatial	10	Project Management	10
Support	10	Testing	10
Product	9	Paid Media	7
Game Development	5	Spatial Computing	3

LICENSING & PRIVACY

- ▶ Source-available under the AgentForge Non-Commercial Source-Available License (ANCSAL) v1.0 — free to use, run, and modify for personal and non-commercial purposes; commercial licensing available on request.
- ▶ No analytics, no tracking, no first-party backend — the only network calls are the AI provider you configure, optional web search, optional code execution, and public agent-prompt fetches.
- ▶ API keys and all application state are stored exclusively in the browser's `localStorage`.

Ready to put **247 experts** to work?

AgentForge is live, free, and open right now — no signup, no install required to start. Open it on any phone or desktop and pick a specialist.

mmuzammul.github.io/AgentForge

Source & full agent roster: github.com/mmuzammul/AgentForge